

Creating Other Schema Objects

Objectives

- After completing this lesson, you should be able to do the following:
 - Create simple and complex views
 - Retrieve data from views
 - Create, maintain, and use sequences
 - Create and maintain indexes
 - Create private and public synonyms

Database Objects

Object	Description
Table	Basic unit of storage; composed of rows
View	Logically represents subsets of data from one or more tables
Sequence	Generates numeric values
Index	Improves the performance of some queries
Synonym	Gives alternative names to objects

What Is a View?

EMPLOYEES table

EMPLOYEE_ID	FIRST_NAME	LAST_NAME	EMAIL	PHONE_NUMBER	HIRE_DATE	JOB_ID	SALA
100	Steven	King	SKING	515.123.4567	17-JUN-87	AD_FRES	2401
101	Neena	Kochhar	NKOCHHAR	515.123.4568	21-SEP-89	AD_VP	1701
102	Lex	De Haan	LDEHAAN	515.123.4569	13-JAN-93	AD_VP	1701
103	Alexander	Hunold	AHUNOLD	590.423.4567	03-JAN-90	IT_PROG	901
104	Bruce	Ernst	BERNST	590.423.4568	21-MAY-91	IT_PROG	601
107	Diana	Lorentz	DLORENTZ	590.423.5567	07-FEB-98	IT_PROG	421
124	Kevin	Mourgos	KMOURGOS	650.123.5234	16-NOV-99	ST_MAN	581
141	Trenna	Rae	TRAE	650.121.8009	17-OCT-95	ST_CLERK	351
142	Curtis	Denes	CDAVIES	650.121.2994	29-JAN-97	ST_CLERK	311
143	Randall	Mateo	RMATEO	650.121.2074	10-MAR-98	ST_CLERK	261
149	Zlotkey			10500	JUL-96	ST_CLERK	251
174	Abel			11000	JAN-00	SA_MAN	1051
176	Taylor			0600	MAY-96	SA_REP	1101
177	Kimbberly	Grant	KGRANT	611.44.1044.423203	MAR-98	SA_REP	861
178	Kimbberly	Grant	KGRANT	611.44.1044.423203	24-MAY-99	SA_REP	701
200	Jennifer	Whalen	JWHALEN	515.123.4444	17-SEP-87	AD_ASST	441
201	Michael	Hartstein	MHARTSTE	515.123.5555	17-FEB-96	MK_MAN	1301
202	Pat	Fay	PFAY	603.123.6666	17-AUG-97	MK_REP	601
205	Shelley	Higgins	SHIGGINS	515.123.8080	07-JUN-94	AC_MGR	1201
206	William	Gietz	WGIEZT	515.123.8181	07-JUN-94	AC_ACCOUNT	831

20 rows selected.

Advantages of Views

**To restrict
data access**

**To make complex
queries easy**



**To provide
data
independence**

**To present
different views of
the same data**

Simple Views and Complex Views

Feature	Simple Views	Complex Views
Number of tables	One	One or more
Contain functions	No	Yes
Contain groups of data	No	Yes
DML operations through a view	Yes	Not always

Creating a View

- You embed a subquery in the CREATE VIEW statement:

```
CREATE [OR REPLACE] [FORCE|NOFORCE] VIEW view  
  [(alias[, alias]...)]  
  AS subquery  
[WITH CHECK OPTION [CONSTRAINT constraint]]  
[WITH READ ONLY [CONSTRAINT constraint]];
```

Creating a View

- Create the EMPVU80 view, which contains details of employees in department 80:

```
CREATE VIEW empvu80
AS SELECT employee_id, last_name, salary
FROM employees
WHERE department_id = 80;
```

View created.

- Describe the structure of the view by using the *SQL*Plus* DESCRIBE command:

```
DESCRIBE empvu80
```


Creating a View

- Create a view by using column aliases in the subquery:

```
CREATE VIEW    salvu50
  AS SELECT    employee_id ID_NUMBER, last_name NAME,
              salary*12 ANN_SALARY
    FROM      employees
    WHERE     department_id = 50;
View created.
```

Retrieving Data from a View

```
SELECT *  
FROM salvu50;
```

ID_NUMBER	NAME	ANN_SALARY
124	Mourgos	69600
141	Rajs	42000
142	Davies	37200
143	Matos	31200
144	Vargas	30000

Modifying a View

- Modify the EMPVU80 view by using a CREATE OR REPLACE VIEW clause. Add an alias for each column name:

```
CREATE OR REPLACE VIEW empvu80
  (id_number, name, sal, department_id)
AS SELECT  employee_id, first_name || ' '
           || last_name, salary, department_id
  FROM      employees
 WHERE     department_id = 80;
View created.
```

are listed in the same order as the columns in the subquery.



Creating a Complex View

- Create a complex view that contains group functions to display values from two tables:

```
CREATE OR REPLACE VIEW dept_sum_vu
  (name, minsal, maxsal, avgsal)
AS SELECT    d.department_name, MIN(e.salary),
             MAX(e.salary),AVG(e.salary)
  FROM      employees e JOIN departments d
  ON        (e.department_id = d.department_id)
  GROUP BY d.department_name;
```

View created.

Rules for Performing DML Operations on a View

- You can usually perform DML operations on simple views. 
- You cannot remove a row if the view contains the following:
 - Group functions
 - A `GROUP BY` clause
 - The `DISTINCT` keyword
 - The pseudocolumn `ROWNUM` keyword

Rules for Performing DML Operations on a View

- You cannot modify data in a view if it contains:
 - Group functions
 - A `GROUP BY` clause
 - The `DISTINCT` keyword
 - The pseudocolumn `ROWNUM` keyword
 - Columns defined by expressions

Rules for Performing DML Operations on a View

- You cannot add data through a view if the view includes:
 - Group functions
 - A GROUP BY clause
 - The DISTINCT keyword
 - The pseudocolumn ROWNUM keyword
 - Columns defined by expressions
 - NOT NULL columns in the base tables that are not selected by the view

Using the WITH CHECK OPTION Clause

- You can ensure that DML operations performed on the view stay in the domain of the view by using the WITH CHECK OPTION clause:

```
CREATE OR REPLACE VIEW empvu20
AS SELECT      *
   FROM      employees
   WHERE      department_id = 20
   WITH CHECK OPTION CONSTRAINT empvu20_ck ;
```

View created.

- Any attempt to change the department number for any row in the view fails because it violates the WITH CHECK OPTION constraint.

Denying DML Operations

- You can ensure that no DML operations occur by adding the `WITH READ ONLY` option to your view definition.
- Any attempt to perform a DML operation on any row in the view results in an Oracle server error.



Denying DML Operations

```
CREATE OR REPLACE VIEW empvu10
    (employee_number, employee_name, job_title)
AS SELECT      employee_id, last_name, job_id
    FROM        employees
    WHERE       department_id = 10
    WITH READ ONLY ;
```

View created.

Removing a View

- You can remove a view without losing data because a view is based on underlying tables in the database.

```
DROP VIEW view;
```

```
DROP VIEW empvu80;  
View dropped.
```

Practice 10: Overview of Part 1


- This practice covers the following topics:
 - Creating a simple view
 - Creating a complex view
 - Creating a view with a check constraint
 - Attempting to modify data in the view
 - Removing views

Sequences

Object	Description
Table	Basic unit of storage; composed of rows
View	Logically represents subsets of data from one or more tables
Sequence	Generates numeric values
Index	Improves the performance of some queries
Synonym	Gives alternative names to objects

Sequences

- A sequence:
 - Can automatically generate unique numbers
 - Is a sharable object
 - Can be used to create a primary key value
 - Replaces application code
 - Speeds up the efficiency of accessing sequence values when cached in memory



CREATE SEQUENCE

Statement:

Syntax

- Define a sequence to generate sequential numbers automatically:

```
CREATE SEQUENCE sequence
    [INCREMENT BY n]
    [START WITH n]
    [{MAXVALUE n | NOMAXVALUE}]
    [{MINVALUE n | NOMINVALUE}]
    [{CYCLE | NOCYCLE}]
    [{CACHE n | NOCACHE}];
```

Creating a Sequence

- Create a sequence named DEPT_DEPTID_SEQ to be used for the primary key of the DEPARTMENTS table.
- Do not use the CYCLE option.

```
CREATE SEQUENCE dept_deptid_seq  
    INCREMENT BY 10  
    START WITH 120  
    MAXVALUE 9999  
    NOCACHE  
    NOCYCLE;
```

Sequence created.

NEXTVAL and CURRVAL

Pseudocolumns

- `NEXTVAL` returns the next available sequence value. It returns a unique value every time it is referenced, even for different users.
- `CURRVAL` obtains the current sequence value.
- `NEXTVAL` must be issued for that sequence before `CURRVAL` contains a value.

Using a Sequence

- Insert a new department named “Support” in location ID 2500:

```
INSERT INTO departments (department_id,  
                        department_name, location_id)  
VALUES (dept_deptid_seq.NEXTVAL,  
      'Support', 2500);
```

1 row created.

sequence:

```
SELECT dept_deptid_seq.CURRVAL  
FROM dual;
```

Caching Sequence Values

- Caching sequence values in memory gives faster access to those values.
- Gaps in sequence values can occur when:
 - A rollback occurs
 - The system crashes
 - A sequence is used in another table

Modifying a Sequence

- Change the increment value, maximum value, minimum value, cycle option, or cache option:

```
ALTER SEQUENCE dept_deptid_seq  
        INCREMENT BY 20  
        MAXVALUE 999999  
        NOCACHE  
        NOCYCLE;
```

Sequence altered.

Guidelines for Modifying a Sequence

- You must be the owner or have the `ALTER` privilege for the sequence.
- Only future sequence numbers are affected.
- The sequence must be dropped and re-created to restart the sequence at a different number.
- Some validation is performed.
- To remove a sequence, use the `DROP` statement:

```
DROP SEQUENCE dept_deptid_seq;  
Sequence dropped.
```

Indexes

Object	Description
Table	Basic unit of storage; composed of rows
View	Logically represents subsets of data from one or more tables
Sequence	Generates numeric values
Index	Improves the performance of some queries
Synonym	Gives alternative names to objects

Indexes

- An index:
 - Is a schema object
 - Can be used by the Oracle server to speed up the retrieval of rows by using a pointer
 - Can reduce disk I/O by using a rapid path access method to locate data quickly
 - Is independent of the table that it indexes
 - Is used and maintained automatically by the Oracle server



How Are Indexes Created?

- Automatically: A unique index is created automatically when you define a `PRIMARY KEY` or `UNIQUE` constraint in a table definition.



- Manually: Users can create nonunique indexes on columns to speed up access to the rows.



Creating an Index

- Create an index on one or more columns:

```
CREATE INDEX index  
ON table (column[, column]...);
```

- Improve the speed of query access to the LAST_NAME column in the EMPLOYEES table:

```
CREATE INDEX emp_last_name_idx  
ON          employees(last_name);  
Index created.
```

Index Creation Guidelines

Create an index when:	
✓	A column contains a wide range of values
✓	A column contains a large number of null values
✓	One or more columns are frequently used together in a WHERE clause or a join condition
✓	The table is large and most queries are expected to retrieve less than 2% to 4% of the rows in the table
Do not create an index when:	
✗	The columns are not often used as a condition in the query
✗	The table is small or most queries are expected to retrieve more than 2% to 4% of the rows in the table
✗	The table is updated frequently
✗	The indexed columns are referenced as part of an expression

Removing an Index

- Remove an index from the data dictionary by using the DROP INDEX command:

```
DROP INDEX index;
```

- Remove the UPPER_LAST_NAME_IDX index from the data dictionary:

```
DROP INDEX emp_last_name_idx;  
Index dropped.
```

- To drop an index, you must be the owner of the index or have the DROP ANY INDEX privilege.

Synonyms

Object	Description
Table	Basic unit of storage; composed of rows
View	Logically represents subsets of data from one or more tables
Sequence	Generates numeric values
Index	Improves the performance of some queries
Synonym	Gives alternative names to objects

Synonyms

- Simplify access to objects by creating a synonym (another name for an object). With synonyms, you can:
 - Create an easier reference to a table that is owned by another user
 - Shorten lengthy object names

```
CREATE [PUBLIC] SYNONYM synonym  
FOR      object;
```

Creating and Removing Synonyms

- Create a shortened name for the DEPT_SUM_VU view:

```
CREATE SYNONYM d_sum  
FOR dept_sum_vu;  
Synonym Created.
```

```
DROP SYNONYM d_sum;  
Synonym dropped.
```

Summary

- In this lesson, you should have learned how to:
 - Create, use, and remove views
 - Automatically generate sequence numbers by using a sequence generator
 - Create indexes to improve query retrieval speed
 - Use synonyms to provide alternative names for objects

Practice 10: Overview of Part 2

- This practice covers the following topics:
 - Creating sequences
 - Using sequences
 - Creating nonunique indexes
 - Creating synonyms